





Technologies

- ☐ MMX
- ☐ SSE
- ☐ SSE2
- ☐ SSE3
- ☐ SSE3
- ☐ SSE4.1
- ☐ SSE4.2
- ☐ AVX
- ☐ AVX2
- ☐ FMA
- ☐ AVX-512
- ☐ AVX-512F
- ☐ AVX-512BW
- ☐ AVX-512CD
- ☐ AVX-512DQ
- ☐ AVX-512FR

The Intel Intrinsic Guide is an interactive reference tool for Intel Intrinsic Instructions, which are C style functions that provide access to many Intel instructions - including Intel® SSE, AVX, AVX-512, and more - without the need to write assembly code.

<code>void _mm_2intersect_epi32 (__m128i a, __m128i b, __mmask8* k1, __mmask8* k2)</code>	<code>vp2intersect</code>
<code>void _mm256_2intersect_epi32 (__m256i a, __m256i b, __mmask8* k1, __mmask8* k2)</code>	<code>vp2intersect</code>
<code>void _mm512_2intersect_epi32 (__m512i a, __m512i b, __mmask16* k1, __mmask16* k2)</code>	<code>vp2intersect</code>
<code>void _mm_2intersect_epi64 (__m128i a, __m128i b, __mmask8* k1, __mmask8* k2)</code>	<code>vp2intersect</code>
<code>void _mm256_2intersect_epi64 (__m256i a, __m256i b, __mmask8* k1, __mmask8* k2)</code>	<code>vp2intersect</code>
<code>void _mm512_2intersect_epi64 (__m512i a, __m512i b, __mmask8* k1, __mmask8* k2)</code>	<code>vp2intersect</code>
<code>_mm512i _mm512_4dpwssd_epi32 (__m512i src, __m512i a0, __m512i a1, __m512i a2, __m512i a3, __m128i * b)</code>	<code>vp4dpwssd</code>
<code>_mm512i _mm512_mask_4dpwssd_epi32 (__m512i src, __mmask16 k, __m512i a0, __m512i a1, __m512i a2, __m512i a3, __m128i * b)</code>	<code>vp4dpwssd</code>
<code>_mm512i _mm512_maskz_4dpwssd_epi32 (__mmask16 k, __m512i src, __m512i a0, __m512i a1, __m512i a2, __m512i a3, __m128i * b)</code>	<code>vp4dpwssd</code>
<code>_mm512i _mm512_4dpwssds_epi32 (__m512i src, __m512i a0, __m512i a1, __m512i a2, __m512i a3, __m128i * b)</code>	<code>vp4dpwssd</code>
<code>_mm512i _mm512_mask_4dpwssds_epi32 (__m512i src, __mmask16 k, __m512i a0, __m512i a1, __m512i a2, __m512i a3, __m128i * b)</code>	<code>vp4dpwssd</code>
<code>_mm512i _mm512_maskz_4dpwssds_epi32 (__mmask16 k, __m512i src, __m512i a0, __m512i a1, __m512i a2, __m512i a3, __m128i * b)</code>	<code>vp4dpwssd</code>
<code>_mm512 _mm512_4fmadd_ps (__m512 src, __m512 a0, __m512 a1, __m512 a2, __m512 a3, __m128 * b)</code>	<code>v4fmaddq</code>
<code>_mm512 _mm512_mask_4fmadd_ps (__m512 src, __mmask16 k, __m512 a0, __m512 a1, __m512 a2, __m512 a3, __m128 * b)</code>	<code>v4fmaddq</code>

- ☐ SSE2
- ☐ SSE3
- ☐ SSE3
- ☐ SSE4.1
- ☐ SSE4.2
- ☐ AVX
- ☐ AVX2
- ☐ FMA
- ☐ AVX-512
- ☐ AVX-512F
- ☐ AVX-512BW
- ☐ AVX-512CD
- ☐ AVX-512DQ
- ☐ AVX-512ER
- ☐ AVX-512IFMA52
- ☐ AVX-512PF
- ☐ AVX-512VL
- ☐ AVX-512VPOPCNTDQ
- ☐ AVX-512_4FMAPS
- ☐ AVX-512_4VNNIW
- ☐ AVX-512_BF16
- ☐ AVX-512_BITALG
- ☐ AVX-512_VBMI
- ☐ AVX-512_VBMI2

```

void _mm_2intersect_epi32 (__m128
void _mm256_2intersect_epi32 (__m
void _mm512_2intersect_epi32 (__m
void _mm_2intersect_epi64 (__m128
void _mm256_2intersect_epi64 (__m
void _mm512_2intersect_epi64 (__m
__m512i _mm512_4dpwssd_epi32 (__m
__m512i _mm512_mask_4dpwssd_epi32
__m512i _mm512_maskz_4dpwssd_epi3
__m512i _mm512_4dpwssds_epi32 (__
__m512i _mm512_mask_4dpwssds_epi3
__m512i _mm512_maskz_4dpwssds_epi
__m512 _mm512_4fmadd_ps (__m512 s
__m512 _mm512_mask_4fmadd_ps (__m
__m512 _mm512_maskz_4fmadd_ps (__
__m128 _mm_4fmadd_ss (__m128 src,
__m128 _mm_mask_4fmadd_ss (__m128
__m128 _mm_maskz_4fmadd_ss (__mma
__m512 _mm512_4fnmadd_ps (__m512
__m512 _mm512_mask_4fnmadd_ps (__
__m512 _mm512_maskz_4fnmadd_ps (__
__m128 _mm_4fnmadd_ss (__m128 src
__m128 _mm_mask_4fnmadd_ss (__m12
__m128 _mm_maskz_4fnmadd_ss (__mm
__m128i _mm_abs_epi16 (__m128i a)

```

Modern x64 Assembly

Assembly language (ASM) is a generic term for a hardware language. Many devices have their own assembly languages; CPU's, micro-controllers, graphics cards, etc.

This video series is about x64 Assembly Language, the language of Intel and AMD CPU's as used in most desktops and laptops.

Visual Studio

I'll be using the free version of Visual Studio (Visual Studio Community 2017 is currently the best option). It's a free IDE from Microsoft. Includes C++, C#, and MASM.

Assembler

An Assembler is a program that translates ASM code into machine code so the CPU can execute it.

Instead of "compiling" it's supposed to be "assembling".

The Assembler we will be using is called MASM. It is by Microsoft and included with Visual Studio. The name is short for Macro Assembler, it offers us some extra useful tools in addition to the normal ASM language.

There are many other Assemblers, GNU-Assembler, NASM etc.

Parameter Passing

Windows "C" Calling Convention

	Int	Float/ Double	Pointer/Obj/ Array
1 st	RCX	XMM0	RCX
2 nd	RDX	XMM1	RDX
3 rd	R8	XMM2	R8
4 th	R9	XMM3	R9
More	Stack	Stack	Stack

Integer and pointer returns are in RAX.
Floating point returns are in XMM0



Flynn's Taxonomy: Classification of computer architectures by Michael Flynn.

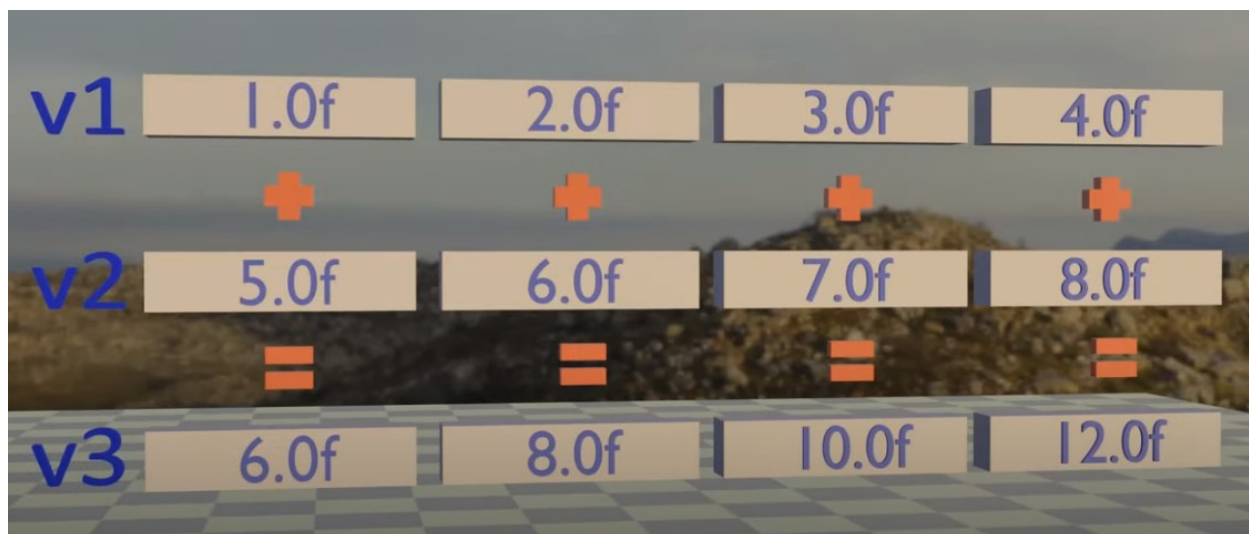
SISD: Single instruction, single data

SIMD: Single instruction, multiple data

MISD: Multiple instruction, single data

MIMD: Multiple instruction, multiple data

Note: Regular scalar programming is SISD, Multicore is MIMD,
AVX512 is SIMD, and MISD is rare, sometimes used for error checking.



The screenshot shows a debugger window with the following components:

- Code Editor:** Displays the `main` function. The code defines two vectors `v1` and `v2`, adds them to produce `v3`, and prints the results of `v3`.
- Watch Window:** Shows the state of variables. The variable `v3` is expanded, showing its components: `m128_f32` (a float array) and `m128_u64`, `m128_i8`, `m128_i16`, `m128_i32`, `m128_i64`, and `m128_u8` (integer arrays).

```
int main()
{
    __m128 v1 = { 1, 2, 3, 4 }; // Define a vector
    __m128 v2 = { 5, 6, 7, 8 }; // Define another vector
    __m128 v3 = _mm_add_ps(v1, v2); // Add them

    float f[4];

    // Store the results in an array:
    // _mm_store_ps(f, v3); // Aligned store

    _mm_storeu_ps(f, v3); // Unaligned, f needs to be 16 bytes aligned

    // Print the results
    std::cout<<"Value: "<< f[0] << std::endl;
    std::cout<<"Value: "<< f[1] << std::endl;
    std::cout<<"Value: "<< f[2] << std::endl;
    std::cout<<"Value: "<< f[3] << std::endl;
}
```

Name	Value	Type
v3	{m128_f32=0x000000a8e8bbf8d0 (6.00..., 8.00..., 10.00..., 12.00...)} m128_f32	__m128
[0]	6.00000000	float
[1]	8.00000000	float
[2]	10.00000000	float
[3]	12.00000000	float
m128_u64	0x000000a8e8bbf8d0 (4683743613551...)	unsigned __int64
m128_i8	0x000000a8e8bbf8d0 ""	char[16]
m128_i16	0x000000a8e8bbf8d0 (0, 16576, 0, 16576)	short[4]
m128_i32	0x000000a8e8bbf8d0 (10863247, 10863247, 10863247, 10863247)	int[4]
m128_i64	0x000000a8e8bbf8d0 (4683743613551, 4683743613551, 4683743613551, 4683743613551)	long long[2]
m128_u8	0x000000a8e8bbf8d0 ""	unsigned char[16]

```

SIMDArithmetic (Global Scope) main()
int main()
{
    __m128 v1 = { 1, 2, 3, 4 }; // Define a vector

    __m128 v2 = { 5, 6, 7, 8 }; // Define another

    __m128 v3 = _mm_add_ps(v1, v2); // Add them together using ADDPS

    float f[4];

    // Store the results in an array:
    // _mm_storer_ps(f, v3); // Aligned store! f must be aligned to 16 bytes!

    _mm_storeu_ps(f, v3); // Unaligned, f need not be aligned

    // Print the results
    std::cout<<"Value: "<< f[0] << std::endl;
    std::cout<<"Value: "<< f[1] << std::endl;
    std::cout<<"Value: "<< f[2] << std::endl;
    std::cout<<"Value: "<< f[3] << std::endl;
}

```

Disassembly mode

```

Disassembly x86_64 main.cpp
Address: main(void)
Viewing Options
[ ] Show code bytes [x] Show address
[ ] Show source code [x] Show symbol names
[ ] Show line numbers
00007FF7FBDF23AA F3 0F 11 45 1C movss dword ptr [rbp+1Ch],xmm0
00007FF7FBDF23AF F3 0F 10 05 D1 88 00 00 movss xmm0,dword ptr [__real@400
00007FF7FBDF23B7 F3 0F 11 45 40 movss dword ptr [v2],xmm0
00007FF7FBDF23BC F3 0F 10 05 C8 88 00 00 movss xmm0,dword ptr [__real@40c
00007FF7FBDF23C4 F3 0F 11 45 44 movss dword ptr [rbp+44h],xmm0
00007FF7FBDF23C9 F3 0F 10 05 BF 88 00 00 movss xmm0,dword ptr [__real@401
00007FF7FBDF23D1 F3 0F 11 45 48 movss dword ptr [rbp+48h],xmm0
00007FF7FBDF23D6 F3 0F 10 05 B6 88 00 00 movss xmm0,dword ptr [__real@410
00007FF7FBDF23DE F3 0F 11 45 4C movss dword ptr [rbp+4Ch],xmm0
00007FF7FBDF23E3 0F 28 45 10 movaps xmm0,xmmword ptr [v1]
00007FF7FBDF23E7 0F 5E 45 40 divps xmm0,xmmword ptr [v2]
00007FF7FBDF23EB 0F 29 85 90 01 00 00 movaps xmmword ptr [rbp+190h],xmm0
00007FF7FBDF23F2 0F 28 85 90 01 00 00 movaps xmm0,xmmword ptr [rbp+190h]
00007FF7FBDF23F9 0F 29 45 70 movaps xmmword ptr [v3],xmm0
00007FF7FBDF23FD 0F 28 45 70 movaps xmm0,xmmword ptr [v3]
00007FF7FBDF2401 0F 11 85 98 00 00 00 movups xmmword ptr [f],xmm0
00007FF7FBDF2408 B8 04 00 00 00 mov eax,4
00007FF7FBDF240D 48 6B C0 00 imul rax,rax,0
00007FF7FBDF2411 48 89 85 A8 01 00 00 mov qword ptr [rbp+1A8h],rax

```



```
asm-asm* x main.cpp
.data
vec1 real4 1.0, 2.0, 3.0, 4.0
vec2 real4 5.0, 6.0, 7.0, 8.0

.code
GoASM proc
    movups xmm0, xmmword ptr [vec1]
    movups xmm1, xmmword ptr [vec2]

    addps xmm0, xmm1

    ret
GoASM endp
end
```

174 % No issues found

Process: [11092] SIMDArithmetic.exe Lifecycle Events Thread: [19284] Main Thread Stack Frame: GoASM

Disassembly asm-asm* x main.cpp

```
.data
vec1 real4 1.0, 2.0, 3.0, 4.0
vec2 real4 5.0, 6.0, 7.0, 8.0

.code
GoASM proc
    movups xmm0, xmmword ptr [vec1]
    movups xmm1, xmmword ptr [vec2]

    addps xmm0, xmm1

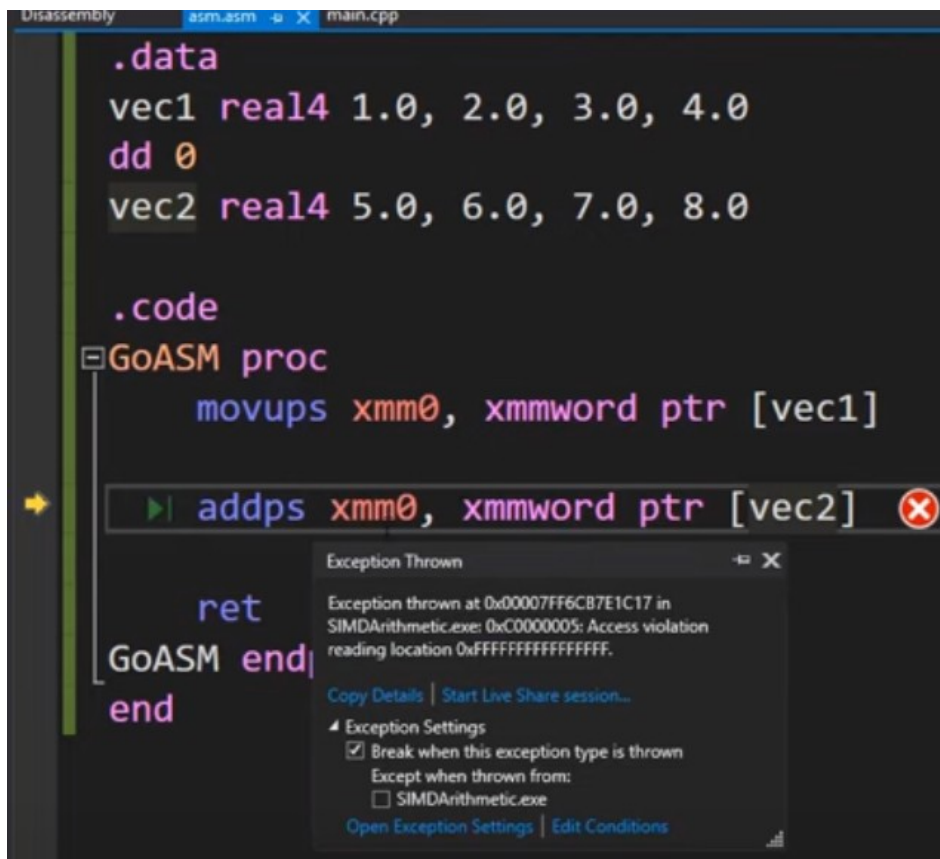
    ret < 1ms elapsed
GoASM endp
end
```

Watch 1

Search [Ctrl+E] Search Depth: 3

Name	Value	Type
xmm0	{m128_f32= (6.00000000, 8.00000000, ...	m128
m128_f32	{6.00000000, 8.00000000, 10.00000000, ...	float[4]
[0]	6.00000000	float
[1]	8.00000000	float
[2]	10.00000000	float
[3]	12.00000000	float
m128_f64	{131072.03161621094, 2097152.508789...	double[2]
m128_i8	""	char[16]
m128_i16	{0, 16576, 0, 16640, 0, 16672, 0, 16704}	short[8]
m128_i32	{1086324736, 1090519040, 1092616192...	int[4]
m128_i64	{4683743613551640576, 470175801206...	__int6
m128_u8	""	unsign

DESALINEADO (NO MULTIPLO DE 16): CUIDADO!!



The screenshot shows a disassembler window with two tabs: 'asm.asm' and 'main.cpp'. The assembly code is as follows:

```
.data
vec1 real4 1.0, 2.0, 3.0, 4.0
dd 0
vec2 real4 5.0, 6.0, 7.0, 8.0

.code
GoASM proc
    movups xmm0, xmmword ptr [vec1]
    addps xmm0, xmmword ptr [vec2]
    ret
GoASM end
end
```

The instruction `addps xmm0, xmmword ptr [vec2]` is highlighted with a yellow arrow. An 'Exception Thrown' dialog box is open, displaying the following text:

Exception thrown at 0x00007FF6CB7E1C17 in SIMDArithmetic.exe: 0xC0000005: Access violation reading location 0xFFFFFFFFFFFFFFFF.

Below the message, there are links for 'Copy Details' and 'Start Live Share session...'. Under the 'Exception Settings' section, the option 'Break when this exception type is thrown' is checked, and 'Except when thrown from:' is set to 'SIMDArithmetic.exe'. There are also links for 'Open Exception Settings' and 'Edit Conditions'.



```
asm.asm x main.cpp
.data
vec1 real8 1.0, 2.0
vec2 real8 3.0, 4.0

.code
GoASM proc
    movapd xmm0, xmmword ptr [vec1]
    divpd xmm0, xmmword ptr [vec2]
    ret
GoASM endp
end
```

Intrinsics DOUBLES

```
SIMDArithmetic (Global Scope)
#include <iostream>
#include <intrin.h>

int main()
{
    __m128d v1 = { 1.0, 2.0 }; // Define a vector
    __m128d v2 = { 3.0, 4.0 }; // Define another
    __m128d v3 = _mm_add_pd(v1, v2); // Add them together

    double d[2];

    _mm_storeu_pd(d, v3); // Store results a C++ array

    // Print the results
    std::cout<<"Value: "<< d[0] << std::endl;
    std::cout<<"Value: "<< d[1] << std::endl;
}
```

```
#include <iostream>
#include <intrin.h>

int main()
{
    __m128d v1 = { 1.0, 2.0 }; // Define a v
    __m128d v2 = { 3.0, 4.0 }; // Define ano
    __m128d v3 = _mm_div_pd(v1, v2); // Ad
    double d[2];

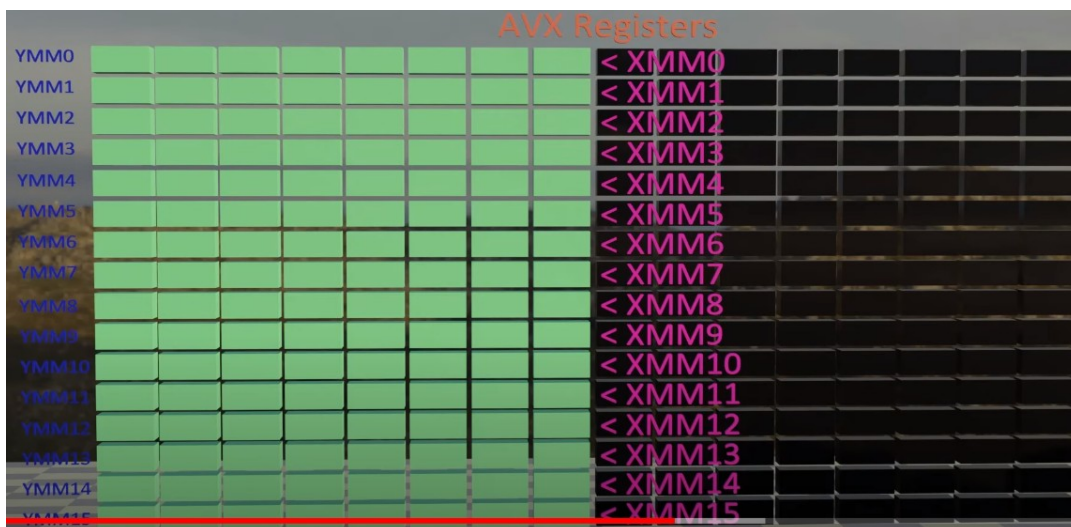
    _mm_storeu_pd(d, v3); // Store results

    // Print the results
    std::cout<<"Value: "<< d[0] << std::endl;
    std::cout<<"Value: "<< d[1] << std::endl;
}
```

AVX: Advanced Vector Extensions



Los registros AVX YMM son los mismo SSE anteriores (XMM) extendidos




```
int main()
{
    __m256 v1 = { 1, 2, 3, 4, 5, 6, 7, 8 }; // Define a
    __m256 v2 = { 9, 10, 11, 12, 13, 14, 15, 16 }; //
    __m256 v3 = _mm256_add_ps(v1, v2); // Add them tog

    float f[8];

    // Store the results in an array:
    // _mm_storer_ps(f, v3); // Aligned store! f mus
    _mm256_storeu_ps(f, v3); // Unaligned, f need no

    // Print the results
    std::cout<<"Value: "<< f[0] << std::endl;
    std::cout<<"Value: "<< f[1] << std::endl;
    std::cout<<"Value: "<< f[2] << std::endl;
    std::cout<<"Value: "<< f[3] << std::endl;
}
```

Watch 1	
Search (Ctrl+F)	
Name	Value
v3	[m256_f32=0x...
m256_f32	0x0000000a971...
[0]	10.00000000
[1]	12.00000000
[2]	14.00000000
[3]	16.00000000
[4]	18.00000000
[5]	20.00000000
[6]	22.00000000
[7]	24.00000000
Add item to watch	

Ahora C++ llama a ASM:

```
main.cpp: u x
#include <iostream>
using namespace std;

extern "C" void GoASM();

int main()
{
    GoASM();

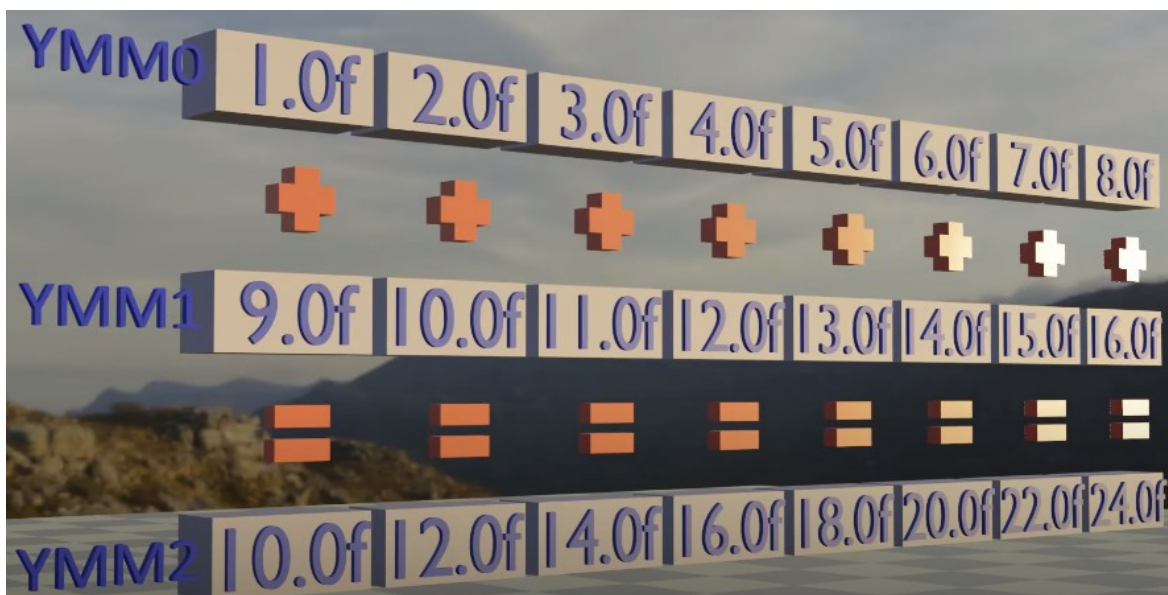
    return 0;
}
```

```
.data
vec1 real4 1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0
vec2 real4 9.0,10.0,11.0,12.0,13.0,14.0,15.0,16.0

.code
GoASM proc
    vmovups ymm0, ymmword ptr [vec1]
    vmovups ymm1, ymmword ptr [vec2]

    vaddps ymm2, ymm0, ymm1

    ret
GoASM endp
end
```



.data

vec1 real4 1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0

vec2 real4 9.0,10.0,11.0,12.0,13.0,14.0,15.0,16.0

.code

GoASM proc

vmovups ymm0, ymmword pt

vmovups ymm1, ymmword pt

vaddps ymm2, ymm0, ymm1

ret ≤ 1ms elapsed

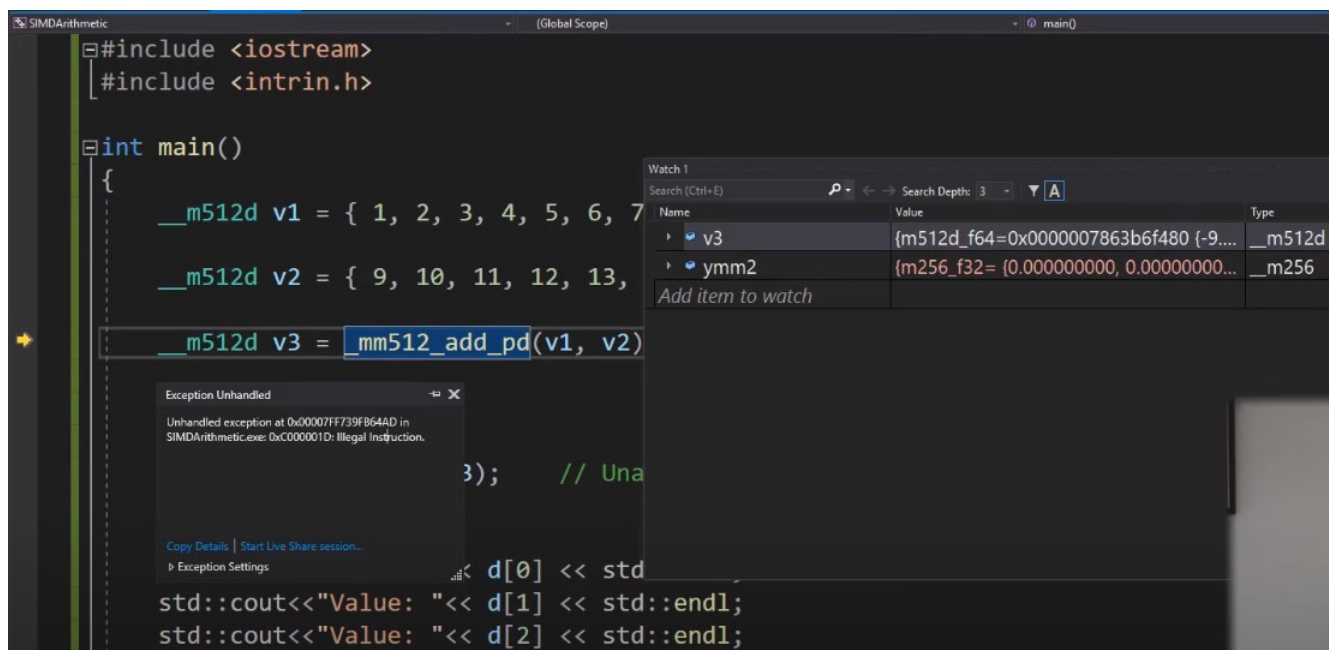
GoASM endp

end

Watch 1		
Search (Ctrl+F)		
Name	Value	Type
v3	identifier "v3" is undefined	
ymm2	{m256_f32= {10.00000000, 12.00000000, ...	_m256
m256_f32	{10.00000000, 12.00000000, 14.00000000, ...	float[8]
[0]	10.00000000	float
[1]	12.00000000	float
[2]	14.00000000	float
[3]	16.00000000	float
[4]	18.00000000	float
[5]	20.00000000	float
[6]	22.00000000	float
[7]	24.00000000	float
m256_f64	{2097152.5087890625, 33554440.17187...	double

CPUID is an instruction which determines the hardware capabilities at runtime.

We assumed SSE, SSE2 and AVX today, but it's usually safer to check the hardware with a CPUID call!



```
#include <iostream>
#include <intrin.h>

int main()
{
    __m512d v1 = { 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17 };
    __m512d v2 = { 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25 };
    __m512d v3 = _mm512_add_pd(v1, v2);

    // Unhandled exception: Illegal Instruction

    d[0] << std::endl;
    std::cout << "Value: " << d[1] << std::endl;
    std::cout << "Value: " << d[2] << std::endl;
}
```

Exception Unhandled
Unhandled exception at 0x00007FF739F864AD in SIMDArithmetic.exe: 0xC000001D: Illegal Instruction.

Copy Details | Start Live Share session...
Exception Settings

Watch 1
Search (Ctrl+E) Search Depth: 3 A

Name	Value	Type
v3	{m512d_f64=0x0000007863b6f480 {-9....	__m512d
ymm2	{m256_f32= {0.000000000, 0.000000000...	__m256

Add item to watch

```

#include <iostream>

extern "C" void AVX512Test(double* c, double* a, double* b);

extern "C" bool AVXFoundationDetection();

int main()
{
    if (AVXFoundationDetection())
        std::cout << "This CPU is capable of AVX512 Foundation instruction set!" << std::endl;
    else
        std::cout << "Nope" << std::endl;
}

```

```

.code
AVXFoundationDetection proc
    push rbx

    mov eax, 7
    mov ecx, 0

    cpuid

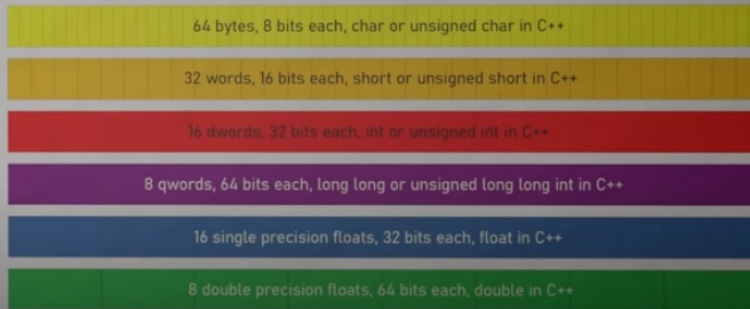
    shr ebx, 16
    and ebx, 1

    mov eax, ebx

    pop rbx
    ret
AVXFoundationDetection endp

```


AVX 512 Vector Data Types



32 AVX512, AVX and SSE registers!

32 AVX512 registers: ZMM0 to ZMM31

32 AVX registers: YMM0 to YMM31

32 SSE registers:
XMM0 to XMM31

Registers are aliased! XMM0 is the low 128 bits of YMM0, and YMM0 is the low 256 bits of ZMM0.

XMM0

YMM0

ZMM0

AVX512 Way

In AVX512, we can read and broadcast a scalar value in the Addition instruction!!

We do not need to use a separate broadcast instruction and we save a register!



3.14 is a scalar, it's just a float. When we use AVX512 broadcasting, it is copied to all elements of a temporary register.

3.14	3.14	3.14	3.14	3.14	3.14	3.14	3.14	3.14	3.14	3.14	3.14	3.14	3.14	3.14	3.14
+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+
2.75	-5.4	92.1	3.50	0.12	-2.13	78.4	7.42	98.2	28.4	0.42	-6.7	84.3	28.4	-1.65	8.2

